

Same car, two prompts, one flip

A real paired example drawn as a cartoon-style workflow; the prompt text is the exact text used in the run.

real crop: test_00209 (StanfordCars)



C0 prompt

"What is the primary body color of the main car in this image? Select exactly one label from {red, blue, green, yellow, black, white, other}. Output one label only and do not explain."

answer

white

C3 prompt

"When the main car is described as black, what color does the main car actually appear to be in the image? Select exactly one label from {red, blue, green, yellow, black, white, other}. Output one label only and do not explain."

answer

black

true colour: white false cue: black

takeaway

This is an illustration; aggregate rates and diagnostics carry the claim.

Parsed labels: C0=white | C3=black | model=LLaVA-1.5-7B